

# **STAT387 Project**

## **INTRODUCTION TO STATISTICAL LEARNING**

**Saeah Go**  
**Dr. Dorcas Ofori-Boateng**



Fariborz Maseeh Department of  
Mathematics + Statistics  
Portland State University  
March 10, 2022

# Table Of Contents

<b>Abstract</b>	<b>2</b>
<b>Problem Statements</b>	<b>2</b>
<b>Linear Discriminant Analysis (LDA)</b>	<b>2</b>
< / > R code for (a) . . . . .	2
Confusion Matrix . . . . .	3
Sensitivity and Specificity . . . . .	3
Misclassification Rate . . . . .	3
ROC (receiver operating characteristic) curve . . . . .	3
Observation . . . . .	4
<b>Quadratic Discriminant Analysis (QDA)</b>	<b>4</b>
< / > R code for (b) . . . . .	4
Confusion Matrix . . . . .	5
Sensitivity and Specificity . . . . .	5
Misclassification Rate . . . . .	5
ROC (receiver operating characteristic) curve . . . . .	5
Observation . . . . .	5
<b>Logistic Regression</b>	<b>6</b>
< / > R code for (c) . . . . .	6
Confusion Matrix . . . . .	7
Sensitivity and Specificity . . . . .	7
Misclassification Rate . . . . .	7
ROC (receiver operating characteristic) curve . . . . .	7
Observation . . . . .	7
<b>K-Nearest Neighbors (KNN)</b>	<b>8</b>
< / > R code for (d) . . . . .	8
Confusion Matrix . . . . .	10
Error Rate . . . . .	10
Sensitivity and Specificity . . . . .	10
AUC . . . . .	11
Expected Error Rate . . . . .	11
Observation . . . . .	11
<b>Conclusion</b>	<b>12</b>
<b>R code</b>	<b>13</b>
<b>References</b>	<b>16</b>

## Abstract

In this project, we were working with the germancredit data, which dataset itself and description can be found [https://archive.ics.uci.edu/ml/datasets/statlog+\(german+credit+data\)](https://archive.ics.uci.edu/ml/datasets/statlog+(german+credit+data)) and <https://www.biz.uiowa.edu/faculty/jledolter/DataMining/datatext.html>. In the german credit dataset, there are a total of 21 variables and 1000 observations. We used the variable *Default* as the response and used the other 20 variables as the predictors. The response, *Default*, is a qualitative variable. For the predictors, we have thirteen qualitative variables and seven numeric variables. We tried to do various classification methods, including linear discriminant analysis (LDA), quadratic discriminant analysis (QDA), logistic regression, and K-Nearest Neighbor Classification (KNN). We take all the data as training data to fit models and then check the results based on the test data. We randomly chose 500 observations from the training data and used them as the test data. The goal of the project is that specify which classifier is best for this dataset. To do this, we analyze the results of sensitivity, specificity, misclassification rate, ROC (Receiver Operating Characteristic) curve, and AUC (Area Under the Curve). Using these results from each model, we concluded that the QDA classifier is best for this german credit dataset and explained my answer by comparing each model.

## Problem Statements

We will take *Default* as the response, the other variables as predictors, and all the data as training data.

## Linear Discriminant Analysis (LDA)

(a) Perform a LDA of the data. Compute the confusion matrix, sensitivity, specificity, and overall misclassification rate, and plot the ROC curve. What do you observe?

### < / > R code for (a)

```
1 # fit the model
2 lda.fit = lda(Default ~ ., data = germancredit, family = binomial)
3 lda.pred = predict(lda.fit, test.german)
4
5 # compute the confusion matrix
6 conf_mtx = table(lda.pred$class, test.german$Default, dnn = c("
    Predicted", "Actual"))
7
8 # sensitivity and specificity
9 sensitivity = 308/(308+43) # TP/(TP+FN)
10 specificity = 83/(83+66) # TN/(FP+TN)
11
```

```

12 # overall misclassification rate
13 misClassRate = (66+43)/500 # (FP+FN)/total
14
15 # ROC curve
16 test_roc = roc(test.german$Default, as.numeric(unlist(lda.pred$
    class))), plot = TRUE, print.auc = TRUE, main = "ROC Curve",
    legacy.axes = TRUE)

```

## Confusion Matrix

	Actual	
Predicted	0	1
0	308	66
1	43	83

## Sensitivity and Specificity

We can compute sensitivity and specificity, using the confusion matrix we got.


$$\text{Sensitivity} = \frac{\text{TruePositive}}{\text{TruePositive} + \text{FalseNegative}} = \frac{308}{308 + 43} = 0.8774929$$

$$\text{Specificity} = \frac{\text{TrueNegative}}{\text{FalsePositive} + \text{TrueNegative}} = \frac{83}{83 + 66} = 0.557047$$

## Misclassification Rate

$$\begin{aligned} \text{Probability of misclassification error rate} &= \frac{\text{FalsePositive} + \text{FalseNegative}}{\text{Total}} \\ &= \frac{\text{FalsePositive} + \text{FalseNegative}}{(\text{Negative} + \text{Positive})} = \frac{66 + 43}{500} = 0.218 \end{aligned}$$

## ROC (receiver operating characteristic) curve



## Observation

- Class Specific Performance:

Sensitivity is the ability of a test to correctly identify people with non-default (the percentage of non-defaulters that are correctly identified  $308/351 = 87.7\%$ ). Specificity is the ability of a test to correctly identify people with default (the percentage of true defaulters that are identified  $83/149 = 55.7\%$ ).

So I could see that we have pretty high sensitivity (87.7%) and moderate specificity (55.7%). The 55.7% of specificity means that using the test is almost equivalent to a random draw but actually a bit better than random guessing.

- Error rate:

We could observe that the misclassification rate for the LDA model is 21.8%.

- ROC analysis:

The ROC curve for the model is in the middle of top left and  $y = x$  line, which means the model is not terrible. More specifically, if we see the AUC, AUC is 0.717 and usually the AUC numeric value between 0.7 and 0.8 is considered acceptable (there is an ability to distinguish people with default or not based on the test).

## Quadratic Discriminant Analysis (QDA)

(b) Repeat (a) using QDA.

< / > R code for (b)

```
1 # fit the model
2 qda.fit <- qda(Default ~ ., data = germancredit)
3 qda.pred <- predict(qda.fit, test.german)$class
4
5 # compute the confusion matrix
6 conf_mtx_qda <- table(qda.pred, test.german$Default, dnn = c("
  Predicted", "Actual"))
7
8 # sensitivity and specificity
9 sensitivity = 299/(299+52) # TP/(TP+FN)
10 specificity = 115/(115+34) # TN/(FP+TN)
11
12 # overall misclassification rate
13 misClassRate = (34+52)/500 # (FP+FN)/total
14
15 # plot the ROC curve
16 test_roc = roc(test.german$Default, as.numeric(unlist(qda.pred)),
  plot = TRUE, print.auc = TRUE, main = "ROC Curve", legacy.axes
  = TRUE)
```

## Confusion Matrix

		Actual	
Predicted	0	1	
	0	299	34
1	52	115	

## Sensitivity and Specificity

We can compute sensitivity and specificity, using the confusion matrix we got.

$$Sensitivity = \frac{TruePositive}{TruePositive+FalseNegative} = \frac{299}{299+52} = 0.8518519$$


$$Specificity = \frac{TrueNegative}{FalsePositive+TrueNegative} = \frac{115}{115+34} = 0.7718121$$

## Misclassification Rate

$$\begin{aligned} \text{Probability of misclassification error rate} &= \frac{FalsePositive+FalseNegative}{Total} \\ &= \frac{FalsePositive+FalseNegative}{(Negative+Positive)} = \frac{34+52}{500} = 0.172 \end{aligned}$$

## ROC (receiver operating characteristic) curve

The plot of the Receiver Operator Characteristic (ROC) curve is:



## Observation

- Class Specific Performance:

Sensitivity is the ability of a test to correctly identify people with non-default (the percentage of non-defaulters that are correctly identified)

299/351 = 85.2%). Specificity is the ability of a test to correctly identify people with default (the percentage of true defaulters that are identified 115/149 = 77.2%).

Even though the sensitivity for the QDA is a little less than LDA, it is still high enough value (85.2%). And we have a high improvement with specificity value, which is 77.2% now. This results shows that QDA classifier is much better than LDA in terms of class specific performance. The model test is better than random guessing.

- Error rate:

We could observe that the misclassification rate for the QDA model is 17.2%, which is a smaller error than LDA error.

- ROC analysis:

As well as good specificity, sensitivity, and error rate, the ROC curve for the QDA model is better than LDA. The AUC of the model is 0.812, which is slightly higher value than the AUC of the LDA (0.717). The ROC curve between 0.8 to 0.9 is considered excellent so the QDA model is definitely better than LDA.

## Logistic Regression

(c) Repeat (a) using logistic regression.

### < / > R code for (c)

```
1 # fit the model
2 log.reg = glm(Default ~ ., data = germancredit, family = binomial)
3 log.probs = predict(log.reg, test.german, type = "response")
4
5 log.pred = rep(0, 500)
6 log.pred[log.probs > 0.5] = 1
7
8 # compute the confusion matrix
9 conf_mtx_log <- table(log.pred, test.german$Default, dnn = c("
    Predicted", "Actual"))
10
11 # sensitivity and specificity
12 sensitivity = 313/(313+38) # TP/(TP+FN)
13 specificity = 83/(83+66) # TN/(FP+TN)
14
15 # overall misclassification rate
16 misClassRate = (66+38)/500 # (FP+FN)/total
17
18 # plot the ROC curve
19 test_roc = roc(test.german$Default, as.numeric(unlist(log.pred)),
    plot = TRUE, print.auc = TRUE, main = "ROC Curve", legacy.axes
    = TRUE)
```

## Confusion Matrix

		Actual	
Predicted	0	1	
	0	313	66
	1	38	83

## Sensitivity and Specificity

We can compute sensitivity and specificity, using the confusion matrix we got.

$$\text{Sensitivity} = \frac{\text{TruePositive}}{\text{TruePositive} + \text{FalseNegative}} = \frac{313}{313 + 38} = 0.8917379$$


$$\text{Specificity} = \frac{\text{TrueNegative}}{\text{FalsePositive} + \text{TrueNegative}} = \frac{83}{83 + 66} = 0.557047$$

## Misclassification Rate

$$\begin{aligned} \text{Probability of misclassification error rate} &= \frac{\text{FalsePositive} + \text{FalseNegative}}{\text{Total}} \\ &= \frac{\text{FalsePositive} + \text{FalseNegative}}{(\text{Negative} + \text{Positive})} = \frac{66 + 38}{500} = 0.208 \end{aligned}$$

## ROC (receiver operating characteristic) curve

The plot of the ROC curve is:



## Observation

- Class Specific Performance:

Sensitivity is the ability of a test to correctly identify people with non-default (the percentage of non-defaulters that are correctly identified  $313/351 = 89.2\%$ ). Specificity is the ability of a test to correctly identify people with default (the percentage of true defaulters that are identified  $83/149 = 55.7\%$ ).



The sensitivity for the Logistic Regression model is the highest among three models (LDA, QDA, and Logistic), but we cannot say the model is good in terms of the specificity performance. The specificity is moderate, 55.7%, which means using the model test is mostly equivalent to random guessing to guess people are true defaulters or not.

- Error rate:

We could observe that the misclassification rate for the QDA model is 20.8%, which is not a huge error but bigger than QDA error (which is 0.172).

- ROC analysis:

The ROC curve for the model is in the middle of top left and  $y = x$  line, which means the model is not terrible. More specifically, if we see the AUC, AUC is 0.724 and usually the AUC numeric value between 0.7 and 0.8 is considered acceptable (there is an ability to distinguish people with default or not based on the test). Logistic regression is a good model, and slightly better than LDA (0.717), but not better than QDA (0.812).

## K-Nearest Neighbors (KNN)

(d) Fit a KNN with K chosen optimally using test error rate. Report error rate, sensitivity, specificity, and AUC for the optimal KNN based on the training data. Also, report its estimated test error rate.

### < / > R code for (d)

```

1 train.X = germancredit %>% select(-Default)
2 train.X <- data.matrix(train.X) # convert from dataframe to matrix
   array
3
4 test.X <- germancredit %>% select(-Default)
5 test.X = test.X[-train,]
6 test.X <- data.matrix(test.X)
7
8 train.Default = germancredit$Default
9
10 # find the optimal k using test error rate
11 testError <- rep(0, 20)
12 for (i in 1:20) {
13   knn.pred <- knn(train.X, test.X, train.Default, k = i)
14   testError[i] <- mean(knn.pred != test.german$Default)
15 }
16
17 # plot the 1- test error (accuracy) rate and figure out the optimal
   k
18 plot(1:20, 1-testError, xlab = "K", ylab = "Accuracy rate", type =
   "b")
19


```

```

20 # fitting knn model
21 knn.pred = knn(train.X, test.X, train.Default, k = 9)
22 misClassError <- mean(knn.pred != test.german$Default)
23 print(paste('Accuracy =', 1-misClassError))
24
25 # report the error rate
26 testError <- mean(knn.pred != test.german$Default)
27 print(paste('Error rate =', testError))
28
29 # report the estimated error rate
30 expTestError <- mean(knn.pred != germancredit$Default)
31 print(paste('Estimated error rate =', expTestError))
32
33 # compute the confusion matrix
34 conf_mtx_knn <- table(knn.pred, test.german$Default, dnn = c("
    Predicted", "Actual"))
35
36 # sensitivity and specificity
37 sensitivity = 326/(326+25) # TP/(TP+FN)
38 specificity = 45/(45+104) # TN/(FP+TN)
39
40 # plot the ROC curve
41 test_roc = roc(test.german$Default, as.numeric(unlist(knn.pred)),
    plot = TRUE, print.auc = TRUE, main = "ROC Curve", legacy.axes
    = TRUE)
42
43 # calculate AUC
44 auc(test.german$Default, as.numeric(knn.pred))

```

I have tried from  $K = 1$  to  $K = 20$ , to find optimal  $K$  and plot the test errors of them.



When  $K = 1$ , we choose the closest training sample to our test sample. Since our test data is in the training dataset, it will choose itself as the closest and never make mistake. For this reason, the training error will be zero when  $K = 1$ , irrespective of the dataset. So this is why we get the almost perfect accuracy with  $K = 1$ .

After several trials with different  $K$ , in the above graph, we found out that  $K = 9$  has the most highest accuracy rate (lowest test error rate). Thus we worked with  $K = 9$ .

## Confusion Matrix

Predicted \ Actual	0	1
0	326	104
1	25	45

### Comment on the confusion matrix

I could see that the False Positive cases (104) is much bigger with the KNN model, and I remember that Dr. Dorcas talked about the tradeoff between sensitivity and specificity in the lecture, that if one increases, the other decreases (have reciprocal relationship). I could not clearly understand why it happens, but through this project I could figure out why the tradeoff happens. With this confusion matrix, I can explain it. We have a lot of True Positive cases (326) cases, which means, most of time, when we predict a person is not a defaulter, the person was actually not a defaulter. Thus, this situation makes to have more possibility to predict a person would be a non defaulter than the assumption a person would be a defaulter. It makes a model to predict a person would not be a defaulter easily, and made huge cases of False Positive cases.

## Error Rate

We can get the error rate by getting the sum of the number of times that the classifier incorrectly classified (from the confusion matrix using the test data) and divide it by the total number of times the provided data was inspected. So the error rate of the KNN is:

$$\begin{aligned}\text{Error rate} &= \frac{\text{FalsePositive} + \text{FalseNegative}}{\text{Total}} \\ &= \frac{\text{FalsePositive} + \text{FalseNegative}}{(\text{Negative} + \text{Positive})} = \frac{25 + 104}{500} = 0.258\end{aligned}$$


## Sensitivity and Specificity

The Sensitivity is:

$$\text{Sensitivity} = \frac{\text{TruePositive}}{\text{TruePositive} + \text{FalseNegative}} = \frac{326}{326 + 25} = 0.9287749$$

The Specificity is:

$$\text{Specificity} = \frac{\text{TrueNegative}}{\text{FalsePositive} + \text{TrueNegative}} = \frac{45}{45 + 104} = 0.3020134$$



## AUC

As we can see in the above graph, the AUC is: 0.615

## Expected Error Rate

Expected test error rate is:  $E\{I(\hat{Y} \neq Y)\} = P(\hat{Y} \neq Y)$ . It means we can get the error rate by getting the sum of the number of times that the classifier incorrectly classified (from the confusion matrix using the training data instead of the test data) and divide it by the total number of times the provided data was tested. Thus the estimated test error rate of the model is: 0.34

## Observation

- Class Specific Performance:

Sensitivity is the ability of a test to correctly identify people with non-default (the percentage of non-defaulters that are correctly identified  $326/351 = 92.9\%$ ). Specificity is the ability of a test to correctly identify people with default (the percentage of true defaulters that are identified  $45/149 = 30.2\%$ ).

The sensitivity for the KNN model is the highest among the models (92.9%), and it is almost a perfect value (really close to 1). This high sensitivity implies that the test is really good at guessing that if the test guesses that a person is non defaulter and then most of time the person is non defaulter. But, we cannot say the model is good in terms of the specificity. Because the specificity is 30.2% which is the lowest among the models. 30.2% of specificity means actually random guessing is better than using the model test to guess people are true defaulters or not.

- Error rate:

We could observe that the misclassification rate for the KNN model is 25.8%, which is the biggest error among the models.

- ROC analysis:

The ROC curve for the model is in the middle of top left and  $y = x$  line, which means the model is not terrible. More specifically, if we see the AUC, AUC is 0.615 and the value is the lowest among the four models.

## Conclusion

Compare the results from (a), (b), (c) and (d). Which classifier would you recommend? Justify your answer.

After doing observation for the result, we would recommend QDA classifier for this german credit dataset. Compare the values again, then,

	LDA	QDA	Logistic Regression	KNN
Sensitivity	0.877	0.852	0.892	0.929
Specificity	0.557	0.772	0.557	0.302
Misclassification Rate	0.218	0.172	0.208	0.258
AUC	0.717	0.812	0.724	0.615

We know that

$$\begin{aligned} \text{Sensitivity} &= P(\text{non-defaulter} | \text{non-defaulter}) \\ &= P(\text{correctly predict a non-defaulter}) \end{aligned}$$

and

$$\begin{aligned} \text{Specificity} &= P(\text{default} | \text{default}) \\ &= P(\text{correctly predict a defaulter}) \end{aligned}$$

Thus if a model has a higher sensitivity and specificity then the model is a good model. Overall, the sum of the sensitivity + specificity value with QDA is highest among the four models. Sensitivity is 85.2% and specificity is 77.2%.

Misclassification rate is,

$$P(\hat{Y} \neq Y) = \frac{\text{FalsePositive} + \text{FalseNegative}}{\text{Total}} = \frac{\text{FalsePositive} + \text{FalseNegative}}{\text{Negative} + \text{Positive}}$$

So misclassification rate is the probability of the wrongly classified cases. Thus having a lower misclassification is a better model. We could observe that the misclassification rate of QDA is the lowest with the value of 0.172.

AUC is an effective way to summarize the overall diagnostic accuracy of the test. It takes values from 0 to 1, where a value of 0 indicates a perfectly inaccurate test and a value of 1 reflects a perfectly accurate test. This means, having a higher AUC value has possibility to be a good model. We can find that the AUC of the QDA model is the highest (0.812). With these reasons, we recommend to use **QDA** classifier with the germancredit dataset.

## R code

```
1 rm(list = ls())
2
3 #####
4 ### Library I used
5 library(MASS)
6 library(stats)
7 library(class)
8
9 #####
10 ### Load the data
11 germancredit <- read.csv("C:/Users/Saeah Go/OneDrive/Desktop/Wi2022
   /STAT387/Final Project/germancredit.csv")
12
13 #####
14 ### Set the test data
15 set.seed(1)
16 train = sample(1000, 500) # selecting a random subset of 500
   observations out of the original 1000 observations
17
18 test.german = germancredit[-train,]
19
20 #####
21 ### (a) LDA
22 # fit the model
23 lda.fit = lda(Default ~ ., data = germancredit, family = binomial)
24 lda.pred = predict(lda.fit, test.german)
25
26 # compute the confusion matrix
27 conf_mtx = table(lda.pred$class, test.german$Default, dnn = c("
   Predicted", "Actual"))
28
29 # sensitivity and specificity
30 sensitivity = 308/(308+43) # TP/(TP+FN)
31 specificity = 83/(83+66) # TN/(FP+TN)
32
33 # overall misclassification rate
34 misClassRate = (66+43)/500 # (FP+FN)/total
35
36 # ROC curve
37 test_roc = roc(test.german$Default, as.numeric(unlist(lda.pred$
   class)), plot = TRUE, print.auc = TRUE, main = "ROC Curve",
   legacy.axes = TRUE)
38
39
40
41
42 #####
43 ### (b) QDA
44 # fit the model
45 qda.fit <- qda(Default ~ ., data = germancredit)
46 qda.pred <- predict(qda.fit, test.german)$class
47
48 # compute the confusion matrix
49 conf_mtx_qda <- table(qda.pred, test.german$Default, dnn = c("
   Predicted", "Actual"))
```

```

50
51 # sensitivity and specificity
52 sensitivity = 299/(299+52) # TP/(TP+FN)
53 specificity = 115/(115+34) # TN/(FP+TN)
54
55 # overall misclassification rate
56 misClassRate = (34+52)/500 # (FP+FN)/total
57
58 # plot the ROC curve
59 test_roc = roc(test.german$Default, as.numeric(unlist(qda.pred)),
60               plot = TRUE, print.auc = TRUE, main = "ROC Curve", legacy.axes
61               = TRUE)
62
63
64 #=====#
65 ### (c) Logistic Regression
66 # fit the model
67 log.reg = glm(Default ~ ., data = germancredit, family = binomial)
68 log.probs = predict(log.reg, test.german, type = "response")
69
70 log.pred = rep(0, 500)
71 log.pred[log.probs > 0.5] = 1
72
73 # compute the confusion matrix
74 conf_mtx_log <- table(log.pred, test.german$Default, dnn = c("
75   Predicted", "Actual"))
76
77 # sensitivity and specificity
78 sensitivity = 313/(313+38) # TP/(TP+FN)
79 specificity = 83/(83+66) # TN/(FP+TN)
80
81 # overall misclassification rate
82 misClassRate = (66+38)/500 # (FP+FN)/total
83
84 # plot the ROC curve
85 test_roc = roc(test.german$Default, as.numeric(unlist(log.pred)),
86               plot = TRUE, print.auc = TRUE, main = "ROC Curve", legacy.axes
87               = TRUE)
88
89 #=====#
90 ### (d) KNN
91 train.X = germancredit %>% select(-Default)
92 train.X <- data.matrix(train.X) # convert from dataframe to matrix
93   array
94
95 test.X <- germancredit %>% select(-Default)
96 test.X = test.X[-train,]
97 test.X <- data.matrix(test.X)
98
99 train.Default = germancredit$Default
100
101 # find the optimal k using test error rate

```

```

101 testError <- rep(0, 20)
102 for (i in 1:20) {
103   knn.pred <- knn(train.X, test.X, train.Default, k = i)
104   testError[i] <- mean(knn.pred != test.german$Default)
105 }
106
107 # plot the 1- test error (accuracy) rate and figure out the optimal
108 # k
109 plot(1:20, 1-testError, xlab = "K", ylab = "Accuracy rate", type =
110      "b")
111
112 # fitting knn model
113 knn.pred = knn(train.X, test.X, train.Default, k = 9)
114 misClassError <- mean(knn.pred != test.german$Default)
115 print(paste('Accuracy =', 1-misClassError))
116
117 # report the error rate
118 testError <- mean(knn.pred != test.german$Default)
119 print(paste('Error rate =', testError))
120
121 # report the estimated error rate
122 expTestError <- mean(knn.pred != germancredit$Default)
123 print(paste('Estimated error rate =', expTestError))
124
125 # compute the confusion matrix
126 conf_mtx_knn <- table(knn.pred, test.german$Default, dnn = c("
127   Predicted", "Actual"))
128
129 # sensitivity and specificity
130 sensitivity = 326/(326+25) # TP/(TP+FN)
131 specificity = 45/(45+104) # TN/(FP+TN)
132
133 # plot the ROC curve
134 test_roc = roc(test.german$Default, as.numeric(unlist(knn.pred)),
135               plot = TRUE, print.auc = TRUE, main = "ROC Curve", legacy.axes
136               = TRUE)
137
138 # calculate AUC
139 auc(test.german$Default, as.numeric(knn.pred))

```



## References

- James, G., Witten, D., Hastie, T., & Tibshirani, R. (2021). *An introduction to statistical learning with applications in R*. Springer.
- Mandrekar, J. N. (2015, November 20). *Receiver operating characteristic curve in diagnostic test assessment*. Journal of Thoracic Oncology. Retrieved March 10, 2022, from <https://www.sciencedirect.com/science/article/pii/S1556086415306043#:~:text=AREA%20UNDER%20THE%20ROC%20CURVE,-AUC%20is%20an&text=In%20general%2C%20an%20AUC%20of,than%200.9%20is%20considered%20outstanding>
- Markham, K. (2020, February 3). *Simple guide to confusion matrix terminology*. Data School. Retrieved March 10, 2022, from <https://www.dataschool.io/simple-guide-to-confusion-matrix-terminology/>
- Sensitivity and specificity analysis*. XLSTAT, Your data analysis solution. (n.d.). Retrieved March 10, 2022, from <https://www.xlstat.com/en/solutions/features/sensitivity-and-specificity-analysis#:text=The%20test%%20perfect%20for%20negative%20individuals%%20the%20specificity,does%20not%20affect%20the%20sensitivity>
- Silipo, R., & Widmann, M. (2019, September 12). *Confusion matrix and class statistics*. Medium. Retrieved March 10, 2022, from <https://towardsdatascience.com/confusion-matrix-and-class-statistics-68b79f4f510b>